

AUTONATEAI · SEASON 1



PROGRAM BRIEF

THE CONTEXT WINDOW

A six-episode AI-literacy and data course for students and early-career engineers

A six-episode animated CS series. Five people at a tiny AI studio have thirty days until Demo Day — and every episode, the way they work with AI breaks in a new, very human way. The Operator, the voice in everybody's head, takes us inside the machine (and the brain) to show why, then teaches the fix with real prompts, real code, and real SQL.

Why this, why now

Every student and early-career engineer now works with AI tools daily — but almost none are taught how to work with them: how to specify, how to give context, how to verify, and how to give a model the right facts from a real system of record. Those four skills separate people who are accelerated by AI from people who are quietly misled by it.

6

animated episodes · 26 minutes total

6

hands-on exercises with templates
students keep

12

written lessons across two modules

WHO IT'S FOR

Deans, CS and data-science faculty, career-services and experiential-learning directors, and L&D leaders at employers who hire them.

LEARNING OUTCOMES

- Write specific, structured prompts (role, task, constraints, format, definition of done) for real engineering tasks
- Engineer context: choose what a model sees, summarize instead of paste, and maintain a project context file
- Verify AI output against ground truth and recognize package hallucinations and slopsquatting risk
- Design a relational schema, query it with SQL, and spot JOIN fan-out errors
- Recognize when data is a graph, traverse it with a recursive query, and use graph retrieval as AI context

WHAT'S IN THE BOX

Video	6 animated episodes, 3.5–5 minutes each (~26 minutes total), 16:9, captioned
Workbook	One exercise per episode with keepable prompt/code templates and sanity checks (6–8 hours of practice)
Source lessons	Two full written modules (Prompt & Context Engineering; Relational Databases & Graphs), 12 lessons
Science	Each episode connects the skill to cognitive science and neural networks, with primary sources cited

The episode guide

#	EPISODE	THE SKILL	LENGTH
1	Garbage In	A vague prompt gets the most average answer. A trusted answer gets you a confident mistake.	4:54
2	The Room, Not Just the Ask	A model can only work with what's in its context window. Too little and it guesses; too much and it drowns.	4:21
3	Role, Task, Constraints, Format	A prompt is a scope of work. Four lines turn a wish into a spec — and you decide what done looks like before you start.	3:31
4	That's Not a Source	Fluent isn't factual. A sentence isn't a source — check what would hurt if it were wrong.	4:18
5	The Missing Ledger	Memory is a storyteller; a database is a witness. Model the truth once, then ask it precise questions.	4:31
6	The Web Under the Table	When relationships repeat, loop, and carry their own facts, they're a graph. The connections are the knowledge — for brains, neural nets, and your AI's context.	4:32



EPISODE 1 • 4:54
Garbage In



EPISODE 2 • 4:21
The Room, Not Just the Ask



EPISODE 3 • 3:31
Role, Task, Constraints, Format



EPISODE 4 • 4:18
That's Not a Source



EPISODE 5 • 4:31
The Missing Ledger



EPISODE 6 • 4:32
The Web Under the Table

Watch the season: <https://www.youtube.com/playlist?list=PLVomZxd75ZEs>

Where it fits

- Intro programming (CS1/CS2) — AI-assisted coding norms in the first weeks
- Databases / data-science courses — schema design, SQL correctness, graphs
- Career readiness and internship prep — the skills employers now assume
- Faculty development — a shared vocabulary for responsible AI use in coursework
- Employer onboarding for new grads

DELIVERY OPTIONS

Flipped classroom	Watch the episode before class; run the exercise together in lab.
Workshop series	Six 60–90 minute sessions, one per episode, with a capstone mini-project.
Self-paced	Students complete episodes and the workbook asynchronously; instructors review the capstone.

ASSESSMENT IDEAS

- Before/after prompt rewrite with a written autopsy (Ep 1–3)
- Verification log: three AI suggestions checked against ground truth (Ep 4)
- Mini-database + the COUNT vs COUNT DISTINCT write-up (Ep 5)
- Capstone: a small graph of a real community + an AI-drafted intro grounded in cited rows (Ep 6)

GROUNDING IN THE SCIENCE

Every episode connects the skill to cognitive science and neural networks — working memory, priming, reconstructive memory, confabulation, Hebbian learning — with primary sources cited in the student workbook. Students leave with a mental model of **why** the practices work, not just a list of tips.

RUN A PILOT

A no-friction pilot: one course section or one cohort, all six episodes plus the workbook, a 20-minute faculty walkthrough, and a short outcomes survey at the end. We customize examples to your campus or company if helpful.

Nathan · nate@autonateai.com · autonateai.com